

# Enhancing NHL Salary Evaluation through Dimensionality Reduction

Raphaël Fontaine  
McGill University  
Montreal, Canada  
raphael.fontaine@mail.mcgill.ca

## Feature Selection

```
In [1]: import common
import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler
```

```
In [2]: df = common.load_dataset()

# Split features and label
X_data, y_data = common.split_dataset(df)

X_columns = X_data.columns

# Split train and test data
X_train, y_train, X_test, y_test = common.split_train_test(X_data, y_data)
```

## Feature Selection

### Remove features with low variance

```
In [3]: scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Calculate the variance
variance = np.var(X_train_scaled, axis=0)
variance_df = pd.DataFrame(variance, index=X_columns, columns=['variance'])

# Display the top 10 features with the highest variance
print("Features with the highest variance")
display(variance_df.sort_values(by='variance', ascending=False).head(20))

# Display the top 10 features with the lowest variance
```

```
print("Features with the lowest variance")  
display(variance_df.sort_values(by='variance', ascending=True).head(20))
```

## Features with the highest variance

	variance
<b>nationality_CAN</b>	0.248886
<b>shootsCatches_R</b>	0.232368
<b>position_D</b>	0.228136
<b>nationality_USA</b>	0.186624
<b>position_L</b>	0.159966
<b>position_R</b>	0.133075
<b>season</b>	0.104679
<b>nationality_SWE</b>	0.093311
<b>I_F_shifts_5on4</b>	0.068052
<b>shifts_5on4</b>	0.068052
<b>I_F_oZoneShiftStarts_5on4</b>	0.067822
<b>I_F_flyShiftEnds_5on4</b>	0.063645
<b>icetime_5on4</b>	0.054445
<b>onIce_xGoalsPercentage_5on4</b>	0.054253
<b>I_F_shifts_4on5</b>	0.054040
<b>shifts_4on5</b>	0.054040
<b>I_F_neutralZoneShiftEnds_4on5</b>	0.051402
<b>faceoffsLost</b>	0.050664
<b>faceoffsLost_5on5</b>	0.049979
<b>I_F_dZoneShiftStarts_4on5</b>	0.049933

## Features with the lowest variance

	<b>variance</b>
<b>nationality_NLD</b>	0.000000
<b>nationality_SVN</b>	0.000558
<b>nationality_BLR</b>	0.001394
<b>office_corsiPercentage_4on5</b>	0.002534
<b>nationality_LVA</b>	0.002785
<b>nationality_NOR</b>	0.002785
<b>office_fenwickPercentage_4on5</b>	0.003308
<b>I_F_reboundGoals_4on5</b>	0.003333
<b>office_corsiPercentage_5on4</b>	0.004083
<b>nationality_FRA</b>	0.004171
<b>penaltyMinutesDrawn_4on5</b>	0.004307
<b>office_fenwickPercentage_5on4</b>	0.004616
<b>nationality_DEU</b>	0.004725
<b>I_F_playStopped_4on5</b>	0.005943
<b>team_SJS</b>	0.006106
<b>I_F_dZoneGiveaways_5on4</b>	0.006157
<b>office_xGoalsPercentage_4on5</b>	0.006251
<b>team_SEA</b>	0.006382
<b>office_xGoalsPercentage_5on4</b>	0.006469
<b>I_F_mediumDangerGoals_5on4</b>	0.006663

```

In [4]: variance_thresholds = [0.01, 0.02, 0.03, 0.04, 0.05, 0.06]
        results = []

        for threshold in variance_thresholds:
            selected_features = np.where(variance > threshold)[0]
            X_train_reduced = X_train[:, selected_features]
            X_test_reduced = X_test[:, selected_features]

            # Standardize the data
            X_train_reduced, X_test_reduced = common.standard_scaler(X_train_reduced,
X_test_reduced)

            # Train and evaluate the model
            results_df, predictions = common.train_and_evaluate(X_train_reduced, y_train, X_test_reduced, y_test)

            # Save the results in another df with the specific threshold
            results_df["Features"] = X_train_reduced.shape[1]
            results_df["Threshold"] = threshold

            results.append(results_df)

        results_df = pd.concat(results)
        results_df.groupby(by="Threshold")[results_df.columns].apply(lambda x: x).drop(
        (columns=["Threshold"]))

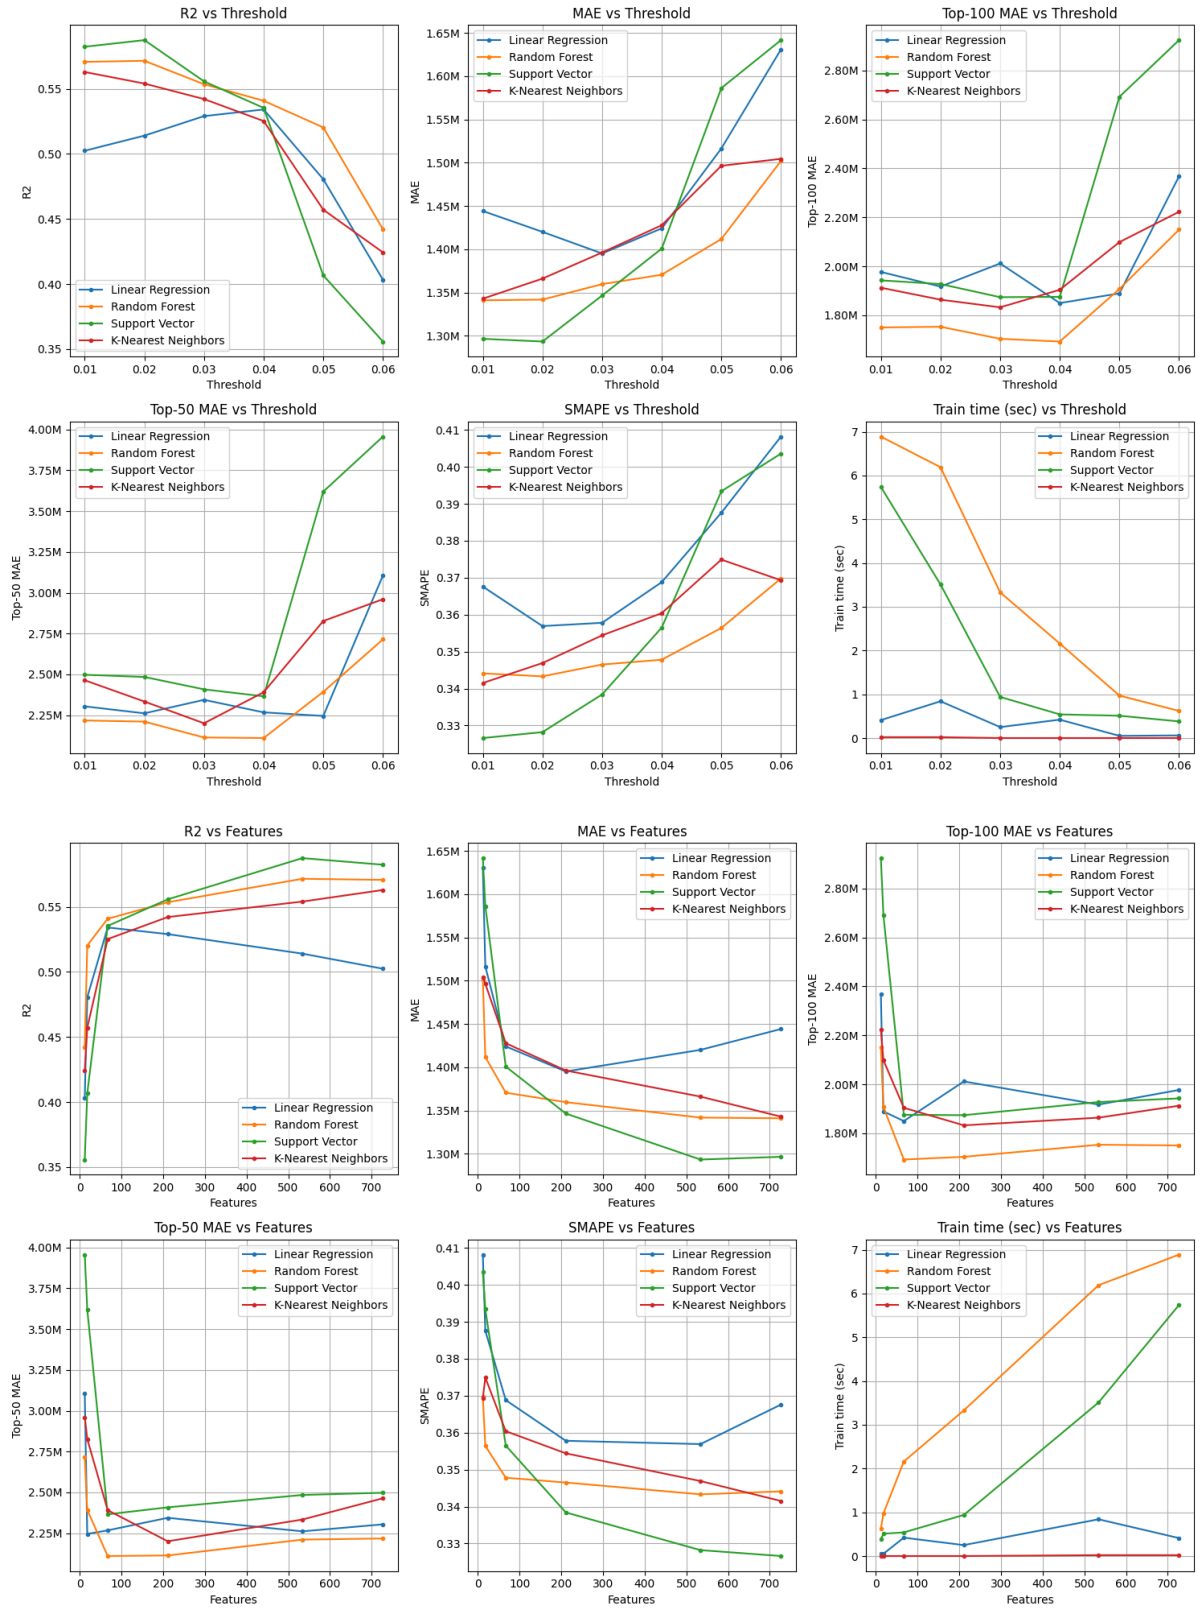
```

Out[4]:

		R2	MAE	Top-100 MAE	Top-50 MAE	SMAPE	Train time (sec)	Features
Threshold	Model							
0.01	Linear Regression	0.5024	1,444,133	1,977,317	2,304,512	0.3676	0.41	727
	Random Forest	0.5707	1,340,970	1,750,731	2,218,032	0.3441	6.89	727
	Support Vector	0.5823	1,296,489	1,942,936	2,497,923	0.3266	5.74	727
	K-Nearest Neighbors	0.5629	1,342,834	1,912,944	2,464,192	0.3415	0.02	727
0.02	Linear Regression	0.5140	1,420,031	1,917,548	2,261,717	0.3569	0.84	534
	Random Forest	0.5715	1,341,772	1,753,626	2,211,082	0.3433	6.19	534
	Support Vector	0.5874	1,293,308	1,927,914	2,484,424	0.3282	3.51	534
	K-Nearest Neighbors	0.5540	1,366,021	1,864,098	2,334,045	0.3469	0.02	534
0.03	Linear Regression	0.5290	1,395,072	2,011,941	2,344,456	0.3578	0.25	211
	Random Forest	0.5535	1,359,619	1,704,305	2,113,901	0.3465	3.33	211
	Support Vector	0.5557	1,346,444	1,874,234	2,408,629	0.3384	0.94	211
	K-Nearest Neighbors	0.5421	1,396,171	1,832,835	2,200,864	0.3544	0.00	211
0.04	Linear Regression	0.5341	1,424,041	1,850,242	2,267,687	0.3688	0.42	67
	Random Forest	0.5408	1,370,534	1,693,126	2,110,504	0.3478	2.16	67
	Support Vector	0.5353	1,400,692	1,875,733	2,365,925	0.3565	0.54	67
	K-Nearest Neighbors	0.5252	1,427,714	1,904,417	2,391,179	0.3604	0.00	67
0.05	Linear Regression	0.4805	1,516,414	1,889,712	2,245,816	0.3876	0.05	18
	Random Forest	0.5203	1,411,781	1,907,684	2,393,466	0.3564	0.97	18
	Support Vector	0.4066	1,586,163	2,692,241	3,619,580	0.3934	0.51	18
	K-Nearest Neighbors	0.4571	1,496,368	2,099,004	2,827,101	0.3749	0.00	18

		R2	MAE	Top-100 MAE	Top-50 MAE	SMAPE	Train time (sec)	Features
Threshold	Model							
0.06	Linear Regression	0.4033	1,630,867	2,367,750	3,106,159	0.4082	0.06	12
	Random Forest	0.4421	1,502,002	2,149,980	2,714,460	0.3698	0.62	12
	Support Vector	0.3558	1,641,583	2,923,764	3,956,213	0.4036	0.38	12
	K-Nearest Neighbors	0.4243	1,504,398	2,222,808	2,960,498	0.3693	0.00	12

```
In [5]: common.plot_metrics(results_df, "Threshold")
common.plot_metrics(results_df, "Features")
```





## Univariate feature selection

```
In [6]: def pearson_correlation(X, y):  
    # Center the data  
    X_centered = X - X.mean()  
    y_centered = y - y.mean()  
  
    # Calculate the covariance  
    covariance = X_centered.T @ y_centered / (X.shape[0] - 1)  
  
    # Calculate the standard deviations  
    std_X = X.std()  
    std_y = y.std()  
  
    # Calculate Pearson's r  
    r = covariance / (std_X * std_y)  
  
    return r
```

```
In [7]: scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Calculate the Pearson correlation
pearson_corr = pearson_correlation(X_train_scaled, y_train)
correlation_df = pd.DataFrame(pearson_corr, index=X_columns, columns=['correlation'])

# Display the top 10 features with the highest correlation
display(correlation_df.sort_values(by='correlation', ascending=False).head(20))

# Display the top 10 features with the lowest correlation
display(correlation_df.sort_values(by='correlation', ascending=True).head(20))
```

	correlation
I_F_oZoneShiftStarts_5on4	0.657052
icetime_5on4	0.624534
Onlce_F_shotAttempts_5on4	0.594658
Onlce_F_scoreAdjustedShotsAttempts_5on4	0.594658
Onlce_F_unblockedShotAttempts_5on4	0.582603
Onlce_F_scoreAdjustedUnblockedShotAttempts_5on4	0.582603
I_F_shifts_5on4	0.581040
shifts_5on4	0.581040
Onlce_F_blockedShotAttempts_5on4	0.577918
Onlce_F_xOnGoal_5on4	0.576197
Onlce_F_shotsOnGoal_5on4	0.575519
Onlce_F_missedShots_5on4	0.574642
Onlce_F_unblockedShotAttempts_other	0.565104
Onlce_F_scoreAdjustedUnblockedShotAttempts_other	0.565104
Onlce_F_lowDangerxGoals_5on4	0.563693
Onlce_F_lowDangerShots_5on4	0.562545
Onlce_A_scoreAdjustedUnblockedShotAttempts_5on4	0.559855
Onlce_A_unblockedShotAttempts_5on4	0.559855
Onlce_F_xGoalsFromxReboundsOfShots_5on4	0.556338
Onlce_F_shotsOnGoal_other	0.553692

	<b>correlation</b>
<b>iceTimeRank</b>	-0.347112
<b>iceTimeRank_5on5</b>	-0.268340
<b>office_corsiPercentage</b>	-0.173838
<b>office_fenwickPercentage</b>	-0.170056
<b>office_xGoalsPercentage</b>	-0.169862
<b>office_corsiPercentage_other</b>	-0.140132
<b>office_fenwickPercentage_other</b>	-0.130943
<b>Office_F_xGoals</b>	-0.113007
<b>I_F_hits_5on5</b>	-0.112420
<b>I_F_hits</b>	-0.104416
<b>nationality_CAN</b>	-0.098920
<b>Office_F_shotAttempts</b>	-0.082560
<b>I_F_dZoneShiftEnds_4on5</b>	-0.081138
<b>position_D</b>	-0.071298
<b>I_F_flyShiftStarts_4on5</b>	-0.069511
<b>OnIce_A_blockedShotAttempts_4on5</b>	-0.069488
<b>OnIce_A_goals_4on5</b>	-0.068182
<b>icetime_4on5</b>	-0.065750
<b>OnIce_A_shotAttempts_4on5</b>	-0.063001
<b>OnIce_A_scoreAdjustedShotsAttempts_4on5</b>	-0.063001

```
In [8]: thresholds = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6]

results = []
for threshold in thresholds:
    selected_features = np.where(np.absolute(pearson_corr) > threshold)[0]
    X_train_reduced = X_train_scaled[:, selected_features]
    X_test_reduced = X_test_scaled[:, selected_features]

    # Standardize the data (scaler)
    X_train_reduced, X_test_reduced = common.standard_scaler(X_train_reduced,
X_test_reduced)

    # Train and evaluate the model
    results_df, predictions = common.train_and_evaluate(X_train_reduced, y_train, X_test_reduced, y_test)

    # Save the results in another df with the specific threshold
    results_df["Features"] = X_train_reduced.shape[1]
    results_df["Threshold"] = threshold

    results.append(results_df)

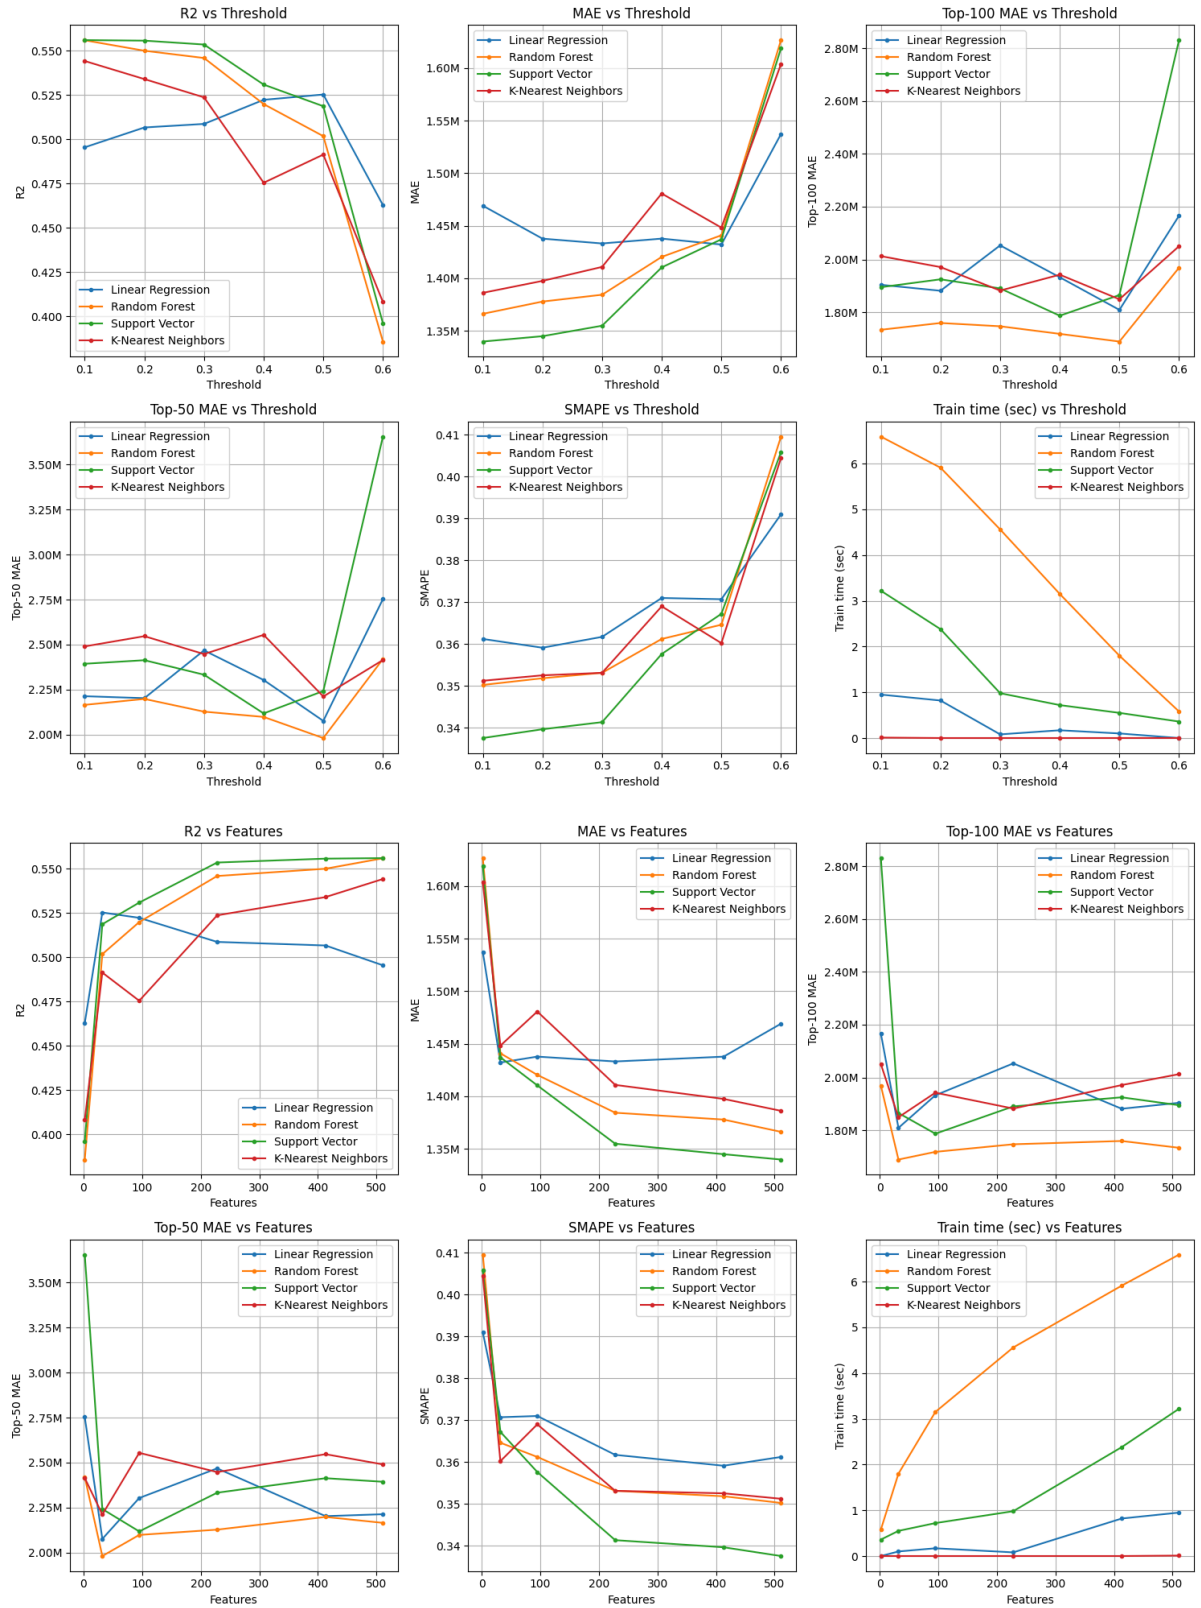
results_df = pd.concat(results)
results_df.groupby(by="Threshold")[results_df.columns].apply(lambda x: x).drop(
(columns=["Threshold"]))
```

Out[8]:

		R2	MAE	Top-100 MAE	Top-50 MAE	SMAPE	Train time (sec)	Features
Threshold	Model							
0.1	Linear Regression	0.4955	1,469,050	1,903,977	2,213,647	0.3612	0.95	511
	Random Forest	0.5559	1,366,071	1,733,658	2,165,628	0.3502	6.59	511
	Support Vector	0.5560	1,339,682	1,894,952	2,393,776	0.3375	3.22	511
	K-Nearest Neighbors	0.5442	1,386,066	2,012,540	2,490,112	0.3512	0.01	511
0.2	Linear Regression	0.5067	1,437,679	1,881,445	2,203,010	0.3591	0.82	413
	Random Forest	0.5500	1,377,752	1,759,376	2,198,645	0.3518	5.91	413
	Support Vector	0.5557	1,344,777	1,924,828	2,413,609	0.3396	2.38	413
	K-Nearest Neighbors	0.5340	1,397,426	1,971,142	2,546,844	0.3525	0.00	413
0.3	Linear Regression	0.5087	1,433,091	2,053,276	2,468,015	0.3617	0.08	228
	Random Forest	0.5459	1,384,257	1,746,963	2,128,031	0.3531	4.56	228
	Support Vector	0.5535	1,354,750	1,890,353	2,333,057	0.3413	0.98	228
	K-Nearest Neighbors	0.5237	1,410,763	1,882,182	2,447,578	0.3531	0.00	228
0.4	Linear Regression	0.5223	1,437,706	1,932,520	2,303,616	0.3710	0.17	95
	Random Forest	0.5199	1,420,296	1,718,210	2,098,734	0.3612	3.15	95
	Support Vector	0.5309	1,410,298	1,786,934	2,118,174	0.3576	0.72	95
	K-Nearest Neighbors	0.4755	1,480,634	1,942,326	2,554,244	0.3690	0.00	95
0.5	Linear Regression	0.5253	1,432,065	1,809,026	2,076,659	0.3707	0.10	32
	Random Forest	0.5018	1,440,867	1,689,103	1,981,560	0.3646	1.80	32
	Support Vector	0.5187	1,437,032	1,865,292	2,242,419	0.3672	0.55	32
	K-Nearest Neighbors	0.4914	1,448,172	1,849,235	2,212,024	0.3602	0.00	32

		R2	MAE	Top-100 MAE	Top-50 MAE	SMAPE	Train time (sec)	Features
Threshold	Model							
0.6	Linear Regression	0.4629	1,537,285	2,165,741	2,752,814	0.3910	0.00	2
	Random Forest	0.3859	1,626,744	1,968,846	2,420,005	0.4096	0.58	2
	Support Vector	0.3962	1,618,978	2,830,753	3,654,712	0.4059	0.36	2
	K-Nearest Neighbors	0.4082	1,603,915	2,050,231	2,414,581	0.4045	0.00	2

```
In [9]: common.plot_metrics(results_df, "Threshold")
common.plot_metrics(results_df, "Features")
```





## Random Selection

```
In [10]: # Standardize the data  
X_train, X_test = common.standard_scaler(X_train, X_test)
```

```
In [11]: n_features = [10, 20, 50, 100, 200, 300, 400, 500, X_data.shape[1]-1]

results = []

for n in n_features:
    # Randomly select n features
    rng = np.random.default_rng(seed=12345)
    selected_features = rng.choice(X_train.shape[1], size=n, replace=False)
    X_train_reduced = X_train[:, selected_features]
    X_test_reduced = X_test[:, selected_features]

    # Train and evaluate the model
    results_df, predictions = common.train_and_evaluate(X_train_reduced, y_train, X_test_reduced, y_test)

    results_df["Features"] = n

    results.append(results_df)

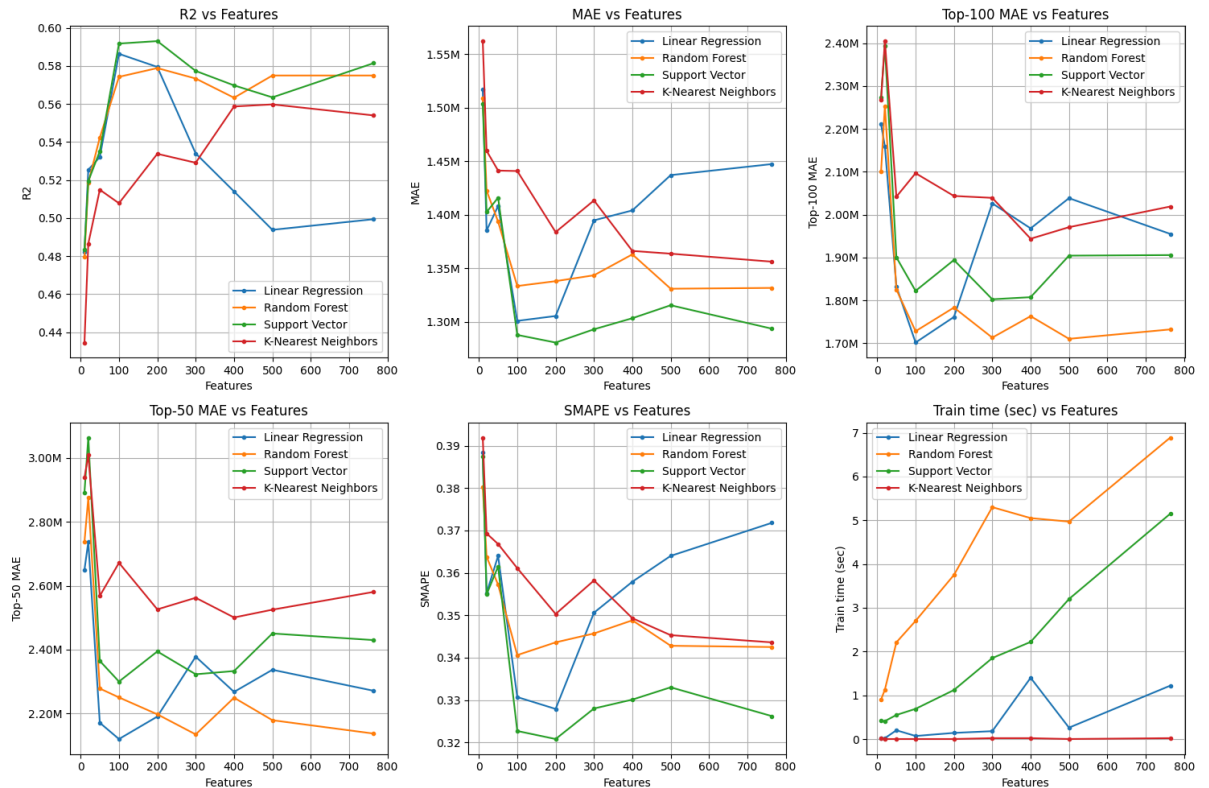
results_df = pd.concat(results)
results_df.groupby(by="Features")[results_df.columns].apply(lambda x: x).drop(
    columns=["Features"])
```

Out[11]:

		R2	MAE	Top-100 MAE	Top-50 MAE	SMAPE	Train time (sec)
Features	Model						
10	Linear Regression	0.4822	1,516,986	2,211,265	2,649,530	0.3884	0.02
	Random Forest	0.4796	1,509,157	2,100,143	2,737,145	0.3803	0.90
	Support Vector	0.4836	1,503,715	2,273,302	2,892,071	0.3875	0.42
	K-Nearest Neighbors	0.4345	1,562,287	2,267,887	2,937,656	0.3919	0.02
20	Linear Regression	0.5254	1,385,284	2,159,127	2,736,993	0.3552	0.02
	Random Forest	0.5186	1,422,653	2,253,787	2,874,604	0.3638	1.13
	Support Vector	0.5194	1,402,562	2,394,130	3,063,242	0.3549	0.41
	K-Nearest Neighbors	0.4865	1,459,765	2,404,914	3,008,550	0.3693	0.00
50	Linear Regression	0.5321	1,408,249	1,832,230	2,171,290	0.3640	0.20
	Random Forest	0.5423	1,393,628	1,825,231	2,278,476	0.3572	2.21
	Support Vector	0.5352	1,416,004	1,900,335	2,364,856	0.3615	0.55
	K-Nearest Neighbors	0.5149	1,441,322	2,041,872	2,568,063	0.3668	0.00
100	Linear Regression	0.5864	1,300,907	1,701,469	2,120,465	0.3307	0.07
	Random Forest	0.5743	1,333,489	1,727,869	2,250,647	0.3406	2.70
	Support Vector	0.5918	1,287,790	1,822,080	2,300,212	0.3227	0.69
	K-Nearest Neighbors	0.5078	1,440,904	2,096,645	2,671,528	0.3611	0.00
200	Linear Regression	0.5795	1,305,477	1,761,106	2,190,760	0.3279	0.14
	Random Forest	0.5789	1,338,019	1,783,135	2,197,795	0.3436	3.75
	Support Vector	0.5931	1,280,671	1,894,051	2,394,359	0.3208	1.12
	K-Nearest Neighbors	0.5338	1,383,844	2,043,542	2,525,929	0.3503	0.00
300	Linear Regression	0.5340	1,394,833	2,026,289	2,378,097	0.3506	0.18
	Random Forest	0.5734	1,343,508	1,712,939	2,134,928	0.3457	5.30
	Support Vector	0.5774	1,293,067	1,802,472	2,323,035	0.3280	1.85
	K-Nearest Neighbors	0.5291	1,413,438	2,039,094	2,562,262	0.3582	0.02
400	Linear Regression	0.5140	1,404,114	1,967,911	2,267,881	0.3579	1.40
	Random Forest	0.5633	1,363,045	1,763,038	2,249,719	0.3488	5.05
	Support Vector	0.5698	1,303,399	1,807,374	2,333,151	0.3301	2.22
	K-Nearest Neighbors	0.5587	1,366,267	1,943,712	2,500,617	0.3493	0.02

		R2	MAE	Top-100 MAE	Top-50 MAE	SMAPE	Train time (sec)
Features	Model						
500	Linear Regression	0.4938	1,437,090	2,038,397	2,337,069	0.3640	0.26
	Random Forest	0.5750	1,330,938	1,710,139	2,179,325	0.3428	4.97
	Support Vector	0.5635	1,315,526	1,904,263	2,450,611	0.3330	3.20
	K-Nearest Neighbors	0.5598	1,363,691	1,970,665	2,525,004	0.3453	0.00
764	Linear Regression	0.4994	1,447,414	1,954,974	2,271,386	0.3718	1.22
	Random Forest	0.5750	1,331,716	1,732,098	2,137,533	0.3425	6.89
	Support Vector	0.5815	1,293,600	1,905,611	2,429,886	0.3262	5.15
	K-Nearest Neighbors	0.5540	1,356,219	2,018,904	2,580,622	0.3436	0.02

In [12]: `common.plot_metrics(results_df, "Features")`



## Curated Features

```
In [13]: ten_features = [
    "age",
    "games_played",
    "icetime",
    "gameScore",
    "position_D",
    "I_F_xGoals",
    "I_F_shotAttempts",
    "I_F_points",
    "I_F_goals",
    "OnIce_A_xGoals"
]

twenty_features = ten_features + [
    "OnIce_A_xGoals",
    "I_F_primaryAssists",
    "I_F_hits",
    "OnIce_F_shotAttempts",
    "I_F_shotsOnGoal",
    "I_F_xOnGoal",
    "OnIce_F_highDangerShots",
    "I_F_xGoals_5on5",
    "I_F_points_5on5",
    "I_F_goals_5on5"
]

thirty_features = twenty_features + [
    "shifts",
    "OnIce_F_xGoals_5on5",
    "onIce_xGoalsPercentage",
    "OnIce_F_xGoals",
    "OnIce_A_xGoals",
    "onIce_corsiPercentage",
    "onIce_fenwickPercentage",
    "I_F_xRebounds",
    "I_F_xPlayContinuedInZone",
    "I_F_rebounds"
]

forty_features = thirty_features + [
    "OnIce_A_highDangerShots",
    "OnIce_A_mediumDangerShots",
    "penalties_5on5",
    "I_F_dZoneGiveaways_5on5",
    "I_F_xGoals_5on4",
    "OnIce_F_goals_5on4",
    "faceoffsWon",
    "OnIce_F_rebounds_5on5",
    "OnIce_F_highDangerShots_5on5",
    "OnIce_F_mediumDangerShots_5on5"
]
```

```

In [14]: results = []

for features in [ten_features, twenty_features, thirty_features, forty_features]:
    selected_features = [X_columns.get_loc(feature) for feature in features]
    X_train_reduced = X_train[:, selected_features]
    X_test_reduced = X_test[:, selected_features]

    # Train and evaluate the models
    results_df, predictions = common.train_and_evaluate(X_train_reduced, y_train, X_test_reduced, y_test)

    results_df["Features"] = len(features)

    results.append(results_df)

results_df = pd.concat(results)
results_df.groupby(by="Features")[results_df.columns].apply(lambda x: x).drop(
    columns=["Features"])

```

Out[14]:

		R2	MAE	Top-100 MAE	Top-50 MAE	SMAPE	Train time (sec)
Features	Model						
10	Linear Regression	0.5972	1,304,619	1,787,843	2,170,968	0.3332	0.02
	Random Forest	0.5480	1,396,831	1,835,874	2,268,770	0.3555	1.00
	Support Vector	0.5421	1,431,640	2,061,207	2,600,320	0.3649	0.35
	K-Nearest Neighbors	0.5208	1,420,575	2,038,903	2,636,208	0.3571	0.02
20	Linear Regression	0.5923	1,306,572	1,790,465	2,139,256	0.3311	0.00
	Random Forest	0.5517	1,402,172	1,758,680	2,235,437	0.3589	1.51
	Support Vector	0.5654	1,382,579	1,808,086	2,252,780	0.3510	0.47
	K-Nearest Neighbors	0.5101	1,455,013	2,010,293	2,527,197	0.3611	0.00
30	Linear Regression	0.5867	1,317,756	1,885,756	2,305,836	0.3354	0.05
	Random Forest	0.5524	1,409,961	1,789,828	2,244,629	0.3603	1.64
	Support Vector	0.5673	1,363,986	1,802,121	2,261,938	0.3434	0.61
	K-Nearest Neighbors	0.5244	1,416,097	1,895,139	2,441,433	0.3515	0.00
40	Linear Regression	0.5846	1,326,821	1,856,544	2,262,506	0.3380	0.06
	Random Forest	0.5581	1,374,543	1,800,319	2,252,110	0.3527	1.83
	Support Vector	0.5750	1,334,249	1,855,225	2,299,091	0.3359	0.55
	K-Nearest Neighbors	0.5262	1,410,998	2,030,788	2,573,531	0.3523	0.00

```
In [15]: common.plot_metrics(results_df, "Features")
```

