

Enhancing NHL Salary Evaluation through Dimensionality Reduction

Raphaël Fontaine
McGill University
Montreal, Canada
raphael.fontaine@mail.mcgill.ca

Partial Least Squares (PLS)

```
In [1]: import common
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cross_decomposition import PLSRegression
```

Load Dataset

```
In [2]: df = common.load_dataset()

# Split features and label
X_data, y_data = common.split_dataset(df)

X_columns = X_data.columns

# Split train and test data
X_train, y_train, X_test, y_test = common.split_train_test(X_data, y_data)

# Standardize the features
X_train, X_test = common.standard_scaler(X_train, X_test)
```

PLS Implementation

```
In [3]: def PLS(n_components):
    X = X_train.astype(np.float64)
    n_samples, n_features = X.shape

    # Reshape y to column vector of floats
    y = y_train.reshape(-1, 1).astype(np.float64)

    # Initialize matrices full of zeros
    T = np.zeros((n_samples, n_components))    # Components
    W = np.zeros((n_features, n_components))    # Weights
    P = np.zeros((n_features, n_components))    # Loadings for X
    Q = np.zeros(n_components)                 # Loadings for y

    for i in range(n_components):
        # Compute weights w that maximize covariance between X and y
        w = X.T @ y
        w /= np.linalg.norm(w)    # Normalize to unit length

        # Project X onto w to find t
        t = X @ w

        tk = (t.T @ t)

        # Compute Loadings p
        p = (X.T @ t) / tk

        # Compute Loadings q
        q = (y.T @ t) / tk
        q = q.item()    # Convert to scalar

        # Deflate X and y
        X -= t @ p.T
        y -= q * t

        # Store results
        T[:, i] = t.ravel()
        P[:, i] = p.ravel()
        W[:, i] = w.ravel()
        Q[i] = q

    return T, W, P, Q
```

```
In [4]: n_components = [5, 10, 25, 50, 100, 200]

results = []
PLS_results = []
for n in n_components:

    # Perform PLS and apply to the data
    T, W, P, Q = PLS(n)
    PLS_results.append((T, W, P, Q))

    X_train_pls = X_train @ W
    X_test_pls = X_test @ W

    # Train and evaluate the models
    results_df, predictions = common.train_and_evaluate(X_train_pls, y_train,
X_test_pls, y_test)

    results_df["Components"] = n

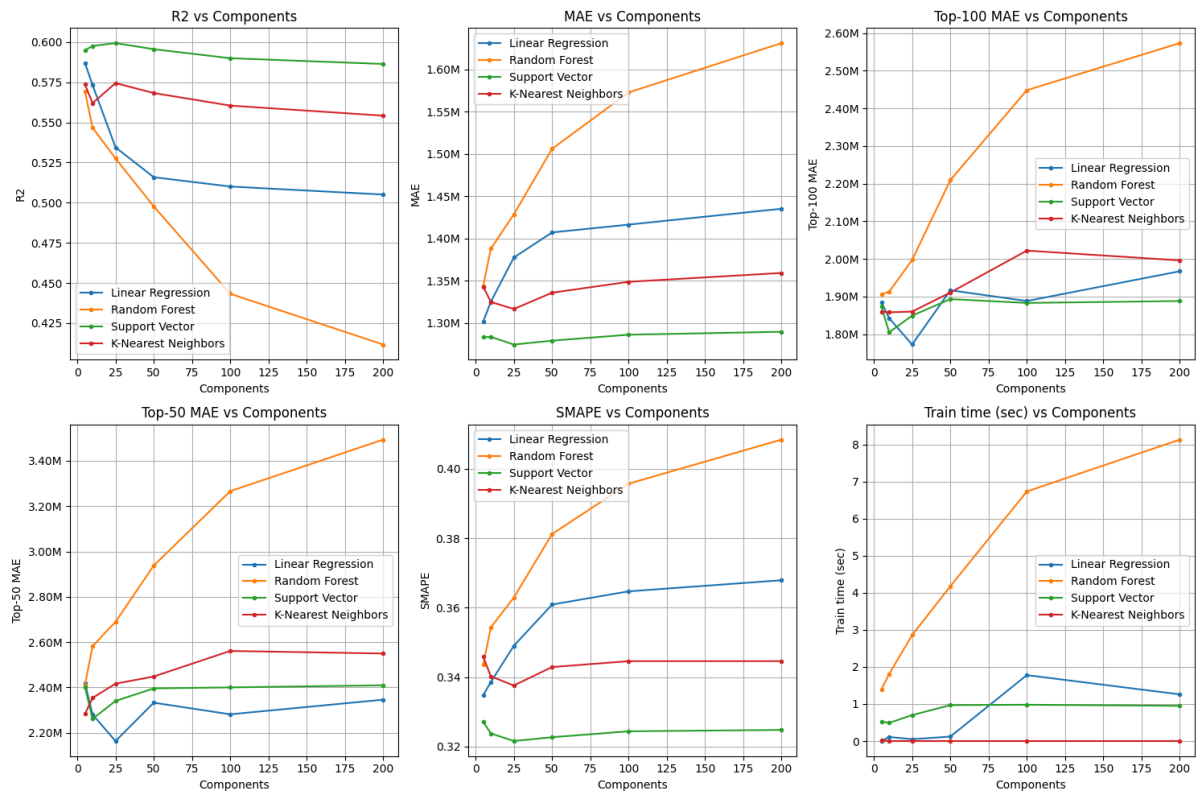
    results.append(results_df)

results_df = pd.concat(results)
results_df.groupby(by="Components")[results_df.columns].apply(lambda x: x).drop
p(columns=["Components"])
```

Out[4]:

		R2	MAE	Top-100 MAE	Top-50 MAE	SMAPE	Train time (sec)
Components	Model						
5	Linear Regression	0.5871	1,301,688	1,884,016	2,420,104	0.3348	0.00
	Random Forest	0.5695	1,344,141	1,906,363	2,408,968	0.3437	1.40
	Support Vector	0.5950	1,283,262	1,873,306	2,402,444	0.3271	0.53
	K-Nearest Neighbors	0.5738	1,343,160	1,860,158	2,282,809	0.3459	0.02
10	Linear Regression	0.5733	1,325,284	1,842,915	2,280,512	0.3386	0.11
	Random Forest	0.5467	1,388,011	1,913,015	2,581,870	0.3543	1.81
	Support Vector	0.5976	1,283,728	1,805,279	2,261,418	0.3238	0.49
	K-Nearest Neighbors	0.5620	1,325,035	1,858,557	2,353,879	0.3402	0.00
25	Linear Regression	0.5345	1,377,533	1,772,794	2,163,478	0.3490	0.05
	Random Forest	0.5276	1,428,605	1,997,287	2,689,512	0.3629	2.86
	Support Vector	0.5994	1,274,635	1,849,491	2,340,399	0.3216	0.70
	K-Nearest Neighbors	0.5745	1,316,685	1,860,251	2,416,610	0.3376	0.00
50	Linear Regression	0.5159	1,407,256	1,916,936	2,332,747	0.3609	0.12
	Random Forest	0.4977	1,506,085	2,209,466	2,938,098	0.3812	4.18
	Support Vector	0.5956	1,279,232	1,893,485	2,395,726	0.3227	0.97
	K-Nearest Neighbors	0.5683	1,335,935	1,910,936	2,448,318	0.3429	0.00
100	Linear Regression	0.5101	1,416,501	1,888,406	2,281,440	0.3647	1.78
	Random Forest	0.4433	1,572,835	2,448,338	3,266,438	0.3957	6.74
	Support Vector	0.5900	1,286,268	1,883,272	2,400,313	0.3244	0.98
	K-Nearest Neighbors	0.5605	1,348,836	2,022,288	2,561,082	0.3446	0.00
200	Linear Regression	0.5051	1,435,270	1,967,415	2,346,021	0.3679	1.26
	Random Forest	0.4117	1,631,209	2,573,281	3,493,959	0.4084	8.14
	Support Vector	0.5864	1,289,766	1,888,387	2,409,532	0.3248	0.95
	K-Nearest Neighbors	0.5542	1,359,316	1,996,245	2,549,803	0.3446	0.00

```
In [5]: common.plot_metrics(results_df, "Components")
```






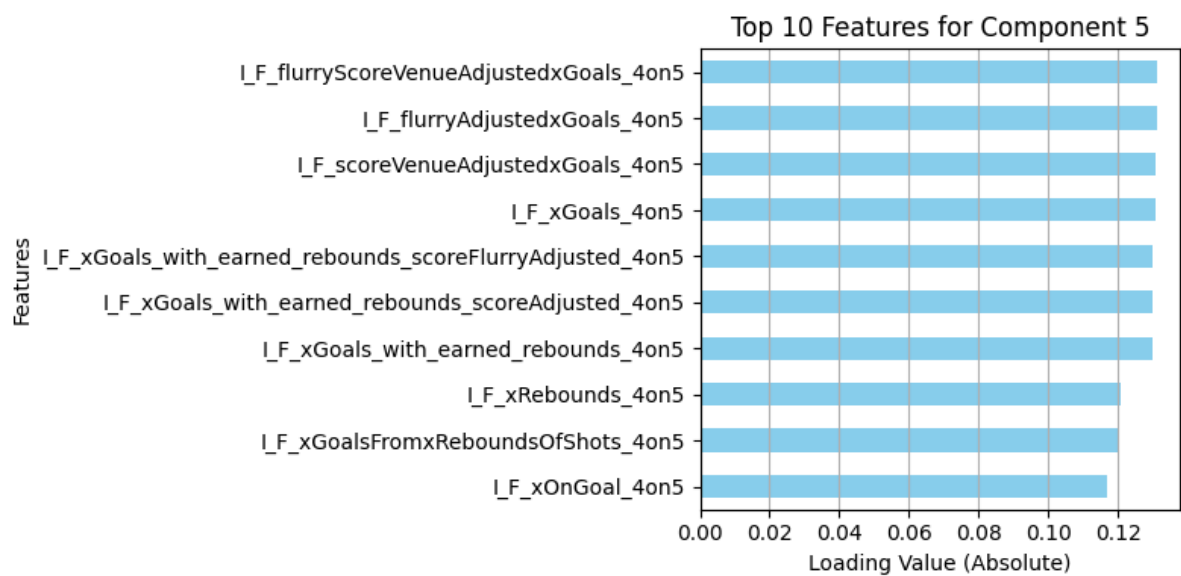

```
In [6]: n_top_features = 10

loadings_df = pd.DataFrame(P, index=X_data.columns)

for comp in range(5):
    plt.figure(figsize=(8, 4))
    # Sort features by absolute loading values for the current component
    top_features = loadings_df.iloc[:, comp].abs().nlargest(n_top_features)

    # Plot
    top_features.sort_values().plot(kind='barh', color='skyblue')
    plt.title(f"Top {n_top_features} Features for Component {comp + 1}")
    plt.xlabel("Loading Value (Absolute)")
    plt.ylabel("Features")
    plt.grid(axis='x')
    plt.tight_layout()
    plt.show()
```





Sklearn PLS

```
In [7]: results = []
        for n in n_components:
            pls = PLSRegression(n_components=n)
            X_train_pls = pls.fit_transform(X_train, y_train)[0]
            X_test_pls = pls.transform(X_test)

            # Train and evaluate the models
            results_df, predictions = common.train_and_evaluate(X_train_pls, y_train,
            X_test_pls, y_test)

            results_df["Components"] = n

            results.append(results_df)

        results_df = pd.concat(results)
        results_df.groupby(by="Components")[results_df.columns].apply(lambda x: x).drop(
        p(columns=["Components"])
```

Out[7]:

		R2	MAE	Top-100 MAE	Top-50 MAE	SMAPE	Train time (sec)
Components	Model						
5	Linear Regression	0.5871	1,301,688	1,884,016	2,420,104	0.3348	0.00
	Random Forest	0.5704	1,315,768	1,992,715	2,591,870	0.3373	1.04
	Support Vector	0.5913	1,291,255	1,909,250	2,449,940	0.3304	0.41
	K-Nearest Neighbors	0.5595	1,336,533	1,888,610	2,421,002	0.3438	0.00
10	Linear Regression	0.5733	1,325,284	1,842,915	2,280,512	0.3386	0.00
	Random Forest	0.5710	1,325,692	1,881,152	2,547,652	0.3376	1.43
	Support Vector	0.5764	1,318,906	1,856,640	2,304,904	0.3357	0.37
	K-Nearest Neighbors	0.5742	1,299,178	1,816,722	2,283,583	0.3311	0.00
25	Linear Regression	0.5345	1,377,533	1,772,794	2,163,478	0.3490	0.00
	Random Forest	0.5631	1,365,039	1,944,955	2,570,970	0.3468	2.40
	Support Vector	0.5471	1,363,322	1,785,004	2,191,026	0.3439	0.65
	K-Nearest Neighbors	0.5745	1,313,526	1,820,760	2,294,476	0.3346	0.00
50	Linear Regression	0.5159	1,407,256	1,916,936	2,332,747	0.3609	0.02
	Random Forest	0.5308	1,435,212	2,150,801	2,850,166	0.3654	3.23
	Support Vector	0.5435	1,371,151	1,854,169	2,272,515	0.3484	0.54
	K-Nearest Neighbors	0.5711	1,323,524	1,814,138	2,353,707	0.3379	0.00
100	Linear Regression	0.5101	1,416,501	1,888,406	2,281,440	0.3647	0.06
	Random Forest	0.4905	1,509,033	2,322,325	3,109,303	0.3836	4.64
	Support Vector	0.5378	1,385,282	1,850,568	2,273,441	0.3522	0.68
	K-Nearest Neighbors	0.5873	1,294,146	1,837,739	2,357,288	0.3292	0.00
200	Linear Regression	0.5051	1,435,270	1,967,415	2,346,021	0.3679	0.07
	Random Forest	0.4440	1,585,455	2,402,792	3,258,788	0.4003	6.79
	Support Vector	0.5346	1,390,313	1,911,683	2,342,793	0.3532	1.06
	K-Nearest Neighbors	0.5817	1,299,766	1,812,803	2,356,020	0.3308	0.00

```
In [8]: common.plot_metrics(results_df, "Components")
```

